

# The Notion of an *Agent* as a Practical Software Engineering Abstraction (Position Statement)

---

Timotheus Kampik, Umeå University, [tkampik@cs.umu.se](mailto:tkampik@cs.umu.se)

In today's interconnected information technology ecosystems, software engineering abstractions that are designed to encapsulate *autonomous* behavior are of increasing relevance. In the academic Artificial Intelligence community, the notion of an *agent* is the most prominent concept in this regard, and an active sub-community works on agent-oriented software engineering abstractions, in particular on *Agent-Oriented Programming* (AOP) languages and frameworks. While existing AOP variants excel at facilitating scientific exploration at the intersection of academic programming and knowledge representation & reasoning, making the case for agent-orientation in practical software engineering is an open problem. To contribute toward a solution, *minimally viable* agent-oriented abstractions can be designed in the context of mainstream programming languages and frameworks. As to our understanding, these abstractions should:

1. Provide a clear and intuitive value proposition to a (perhaps somewhat intellectually curious) practicing software engineer -- even one without a Computer Science degree;
2. Not add any non-essential overhead that steepens the learning curve and hampers adoption, for example by imposing languages, dialects, or paradigms that do not fit into modern software development processes and toolchains;
3. Come with clear design principles with regard to agent internals (reasoning cycle) and agent interface design (interaction and discoverability).

Potential learnings can be drawn from the relatively recent spread of functional programming flavors to mainstream and originally predominantly object-oriented programming languages and frameworks, a high-profile example being React.js in the context of JavaScript (for better or for worse).